

# Chapter\_2\_NNs

January 25, 2022

## 1 Chapter 2: First Neural Network with PyTorch

### Imports

```
[ ]: # importing the necessary libraries

# operating system
import os

# plotting
import matplotlib.pyplot as plt
import seaborn as sns

# linalg
import numpy as np

# dataframes
import pandas as pd

# splitting datasets
from sklearn.model_selection import train_test_split

# metrics - CM + report
from sklearn.metrics import confusion_matrix, classification_report

# neural nets
import torch
import torch.nn as nn
import torch.optim as optim

# functional torch
import torch.nn.functional as F

# colored print statements :D
import termcolor
from termcolor import colored
```

### Loading Data

```
[ ]: # lets load the data located in the data folder
DATA_PATH = '../data/Chapter_2/weatherAUS.csv'
# load into df
df = pd.read_csv(DATA_PATH)
# lets look at the shape of it
#print(df.shape)
# lets look at the columns
#print(df.columns)


# The book suggests keeping these columns:
# cols=['Rainfall', 'Humidity3pm', 'Pressure9am', 'RainToday', 'RainTomorrow']
# I will use the first and the last two only the rest I change
cols =
↳ ['Rainfall', 'WindSpeed9am', 'WindSpeed3pm', 'Pressure3pm', 'RainToday', 'RainTomorrow']
# lets index the df with the respective cols
df = df[cols]
# lets look at the dtypes of the columns
#print(df.dtypes) # we see that 3 are objects, let's look at them

# we see WindDir3PM is about the direction, so let's remove that from the orig
# and replace with windspeed 9 am
# now we can replace the No and Yes with numbers
d = {"No":0, "Yes":1}
df['RainToday'].replace(d, inplace=True)
df['RainTomorrow'].replace(d, inplace=True)

# looking at nanas
#df.isna().sum() # plenty of NaNs, lets drop them

# dropping the nans
df = df.dropna(how='any')

# lets visualize the data
sns.countplot(x = df.RainTomorrow);
plt.show()
# how many of the instances are for rain tomorrow?
print(df.RainTomorrow.value_counts()/df.shape[0]) # high class imbalance
```



```
0.0    0.778723
1.0    0.221277
Name: RainTomorrow, dtype: float64
```

### Splitting Data

```
[ ]: # now lets separate the data
X = df.drop('RainTomorrow',axis=1).values
Y = df.RainTomorrow.values

# lets split the data
TEST_SIZE = 0.2 # 20% for testing, 80% for training
RAND_SEED = 3407
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=TEST_SIZE,
↳random_state=RAND_SEED)
# convert all of them into tensors
# training
X_train = torch.from_numpy(X_train).float()
y_train = torch.from_numpy(y_train).float()
# testing
X_test = torch.from_numpy(X_test).float()
y_test = torch.from_numpy(y_test).float()

# lets look at the training sizes
print(X_train.shape, y_train.shape) # train
```

```
print(X_test.shape, y_test.shape) # test
```

```
torch.Size([100775, 5]) torch.Size([100775])
torch.Size([25194, 5]) torch.Size([25194])
```

## Model

```
[ ]: # lets define a simple neural network
class ANN(nn.Module):

    def __init__(self, n_features):
        super(ANN, self).__init__()
        self.fc1 = nn.Linear(n_features, 15)
        self.fc2 = nn.Linear(15,3)
        self.fc3 = nn.Linear(3,1)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        return torch.sigmoid(self.fc3(x))
```

```
[ ]: # number of features we will have
n_features = X_train.shape[1]
# instantiate the model
model = ANN(n_features)
#lets look at the model
model
```

```
[ ]: ANN(
  (fc1): Linear(in_features=5, out_features=15, bias=True)
  (fc2): Linear(in_features=15, out_features=3, bias=True)
  (fc3): Linear(in_features=3, out_features=1, bias=True)
)
```

## Activation Functions

```
[ ]: # we are using two different activation functions:
# ReLU --> Rectified Linear Unit --> max(0, x)
# Sigmoid --> No specific name --> 1 / (1+e^-x)
# lets look at them
def plot_activation(FUNC,start=-1,stop=1,step=5,title=None, grid=True):
    """
    Helper function to plot the activation functions.
    args:
        start:type:int - defines the start value of the linspace
        stop:type:int - defines the stop value of the linspace
        step:type:int - defines the step of the linspace
        title:type:str - the title to add to the graph
```


```


output:
    plot with the given activation function on the linspace
    '''
    ax = plt.gca()
    space1 = np.linspace(start,stop,step)
    space2 = torch.linspace(start,stop,step)
    # applying the FUNC
    func_space = FUNC(space2).numpy()
    plt.plot(space1, func_space)
    ax.set_ylim([-1.5, 1.5])
    if grid:
        plt.grid()
    if title != None:
        plt.title(title)
    plt.show()

# showing ReLU
plot_activation(F.relu, start=-1, stop=1, step=5,title='ReLU Activation_
↳Function', grid=True)

# showing sigmoid
plot_activation(torch.sigmoid, start=-10, stop=10, step=100,title='Sigmoid_
↳Activation Function',grid=True)

```





**Loss Function and Optimizer** We have: 1. Our Model - Simple ANN with ReLU and Sigmoid  
2. We split our data

What do we need now? 1. Criteria - Loss Function 2. Optimizer - How to we update our weights  
-> Optimizer + Hyperparameters -> Learning Rate

```
[ ]: # define some things
LEARNING_RATE = 0.001
# defining our criteria
criterion = nn.BCELoss() # Binary Cross Entropy -- measures the difference
    ↳ between two binary vectors
# optimizer --> we pass in the models or nets or Neural Networks parameters,
    ↳ and then define the learning rate
opt = optim.Adam(model.parameters(), lr=LEARNING_RATE)
```

### Helper Functions

```
[ ]: # we need to be able to calculate the accuracy of the model

def get_accuracy(y, y_hat):
    """
    Helper function to calc. the accuracy of our models output (y_hat) against
    ↳ the real value y.
    """
    pred = y_hat.ge(0.5).view(-1)
```

```

    return (y == pred).sum().float() / len(y)

# we also need a rounding helper
def round_t(TENSOR, decimal_places=3):
    return round(TENSOR.item(), decimal_places)

```

### Custom Training loop

```

[ ]: def train_loop(MODEL, CRITERIA, OPTIMIZER, EPOCHS, TEST_BATCH):
    '''
        Custom training loop to use the model, criteria and optimizer for the
        ↪ specified number of epochs
    '''
    for EPOCH in range(EPOCHS):
        # TRAINING MODE
        MODEL.train()
        # the prediction of the model
        y_pred = MODEL(X_train)
        y_pred = torch.squeeze(y_pred)
        # training loss
        train_loss = CRITERIA(y_pred, y_train)
        # information from model every nth EPOCH
        if EPOCH%TEST_BATCH == 0:
            # train accuracy
            train_acc = get_accuracy(y_train, y_pred)
            # put the model into EVAL mode!
            MODEL.eval()
            # prediction
            test_pred = MODEL(X_test)
            # squeezing
            test_pred = torch.squeeze(test_pred)
            # calculating our test loss
            test_loss = CRITERIA(test_pred, y_test)
            # test accuracy
            test_acc = get_accuracy(y_test, test_pred)
            # text to print
            train_text = f"TRAINING: EPOCH: {EPOCH} LOSS =
            ↪ {round_t(train_loss)} ACCURACY = {round_t(train_acc)}"
            test_text = f"TESTING: EPOCH: {EPOCH} LOSS = {round_t(test_loss)}
            ↪ ACCURACY = {round_t(test_acc)}"

            print(colored(text=train_text,color='red',on_color='on_grey',attrs=['bold']))

            print(colored(text=test_text,color='green',on_color='on_white',attrs=['bold']))
            # zeroing the gradients - else accumulated gradients from each iter
            OPTIMIZER.zero_grad()
            # Backprop

```

```

train_loss.backward()
# stepping the optimizer
OPTIMIZER.step()

```

```
[ ]: train_lop(model, criterion, opt, 1000,100)
```

```

TRAINING: EPOCH: 0 LOSS = 0.484 ACCURACY = 0.787
TESTING: EPOCH: 0 LOSS = 0.487 ACCURACY = 0.784
TRAINING: EPOCH: 100 LOSS = 0.476 ACCURACY = 0.792
TESTING: EPOCH: 100 LOSS = 0.48 ACCURACY = 0.79
TRAINING: EPOCH: 200 LOSS = 0.476 ACCURACY = 0.792
TESTING: EPOCH: 200 LOSS = 0.48 ACCURACY = 0.79
TRAINING: EPOCH: 300 LOSS = 0.476 ACCURACY = 0.792
TESTING: EPOCH: 300 LOSS = 0.48 ACCURACY = 0.79
TRAINING: EPOCH: 400 LOSS = 0.476 ACCURACY = 0.792
TESTING: EPOCH: 400 LOSS = 0.48 ACCURACY = 0.79
TRAINING: EPOCH: 500 LOSS = 0.476 ACCURACY = 0.792
TESTING: EPOCH: 500 LOSS = 0.48 ACCURACY = 0.79
TRAINING: EPOCH: 600 LOSS = 0.476 ACCURACY = 0.792
TESTING: EPOCH: 600 LOSS = 0.48 ACCURACY = 0.79
TRAINING: EPOCH: 700 LOSS = 0.476 ACCURACY = 0.792
TESTING: EPOCH: 700 LOSS = 0.48 ACCURACY = 0.79
TRAINING: EPOCH: 800 LOSS = 0.476 ACCURACY = 0.792
TESTING: EPOCH: 800 LOSS = 0.48 ACCURACY = 0.79
TRAINING: EPOCH: 900 LOSS = 0.476 ACCURACY = 0.792
TESTING: EPOCH: 900 LOSS = 0.48 ACCURACY = 0.79

```

## Evaluation

```

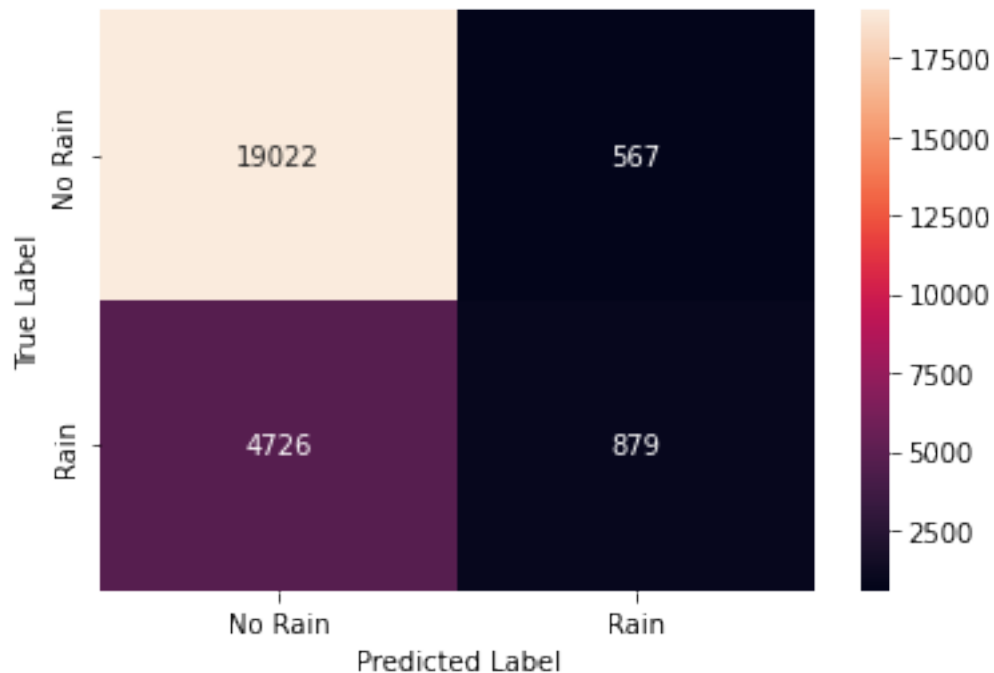
[ ]: # lets look a little deeper into the model
classes = ['No Rain', 'Rain']
# lets predict
predz = model(X_test)
predz = predz.ge(0.5).view(-1)
# classification report
cr = classification_report(y_test, predz, target_names=classes)
print(cr)

```

	precision	recall	f1-score	support
No Rain	0.80	0.97	0.88	19589
Rain	0.61	0.16	0.25	5605
accuracy			0.79	25194
macro avg	0.70	0.56	0.56	25194
weighted avg	0.76	0.79	0.74	25194



```
[ ]: # lets look at a confusion matrix
# format of (tn, fp, fn, tp)
cm = confusion_matrix(y_test, predz)
dfcm = pd.DataFrame(cm, index=classes, columns=classes)
g = sns.heatmap(dfcm,annot=True,fmt='d')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.show()
```



```
[ ]: (tn, fp), (fn, tp) = confusion_matrix(y_test, predz)
print(f"True Negatives: {tn}")
print(f"False Positives: {fp}")
print(f"False Negatives: {fn}")
print(f"True Positives: {tp}")
```

```
True Negatives: 19022
False Positives: 567
False Negatives: 4726
True Positives: 879
```