

Chapter_3_TL

January 27, 2022

1 Chapter 3: Transfer Learning - German Traffic Sign Recognition Benchmark (GTSRB)

Imports

```
[ ]: #files
from glob import glob
import shutil
from pathlib import Path

# linalg
import numpy as np

# images
import cv2

# plotting
from matplotlib.ticker import MaxNLocator
import matplotlib.pyplot as plt
import seaborn as sns
from pylab import rcParams

# dataframes
import pandas as pd

# sklearn
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report

#progress bar
from tqdm import tqdm

# nice colors
import termcolor
from termcolor import colored

# pytorch
import torch, torchvision
```

```

from torch import nn, optim

# functional
import torch.nn.functional as F

# utils
import torchvision.transforms as T
import PIL.Image as Image
from torch.optim import lr_scheduler
from collections import defaultdict

# datasets
from torchvision.datasets import ImageFolder
from torch.utils.data import DataLoader

# models
from torchvision import models

```

```

[ ]: # getting the specific device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
dev_name = torch.cuda.get_device_name()
print(dev_name)

```

Tesla K80

```

[ ]: # setting the figure size
rcParams['figure.figsize'] = 12,8

```

Getting the GTSRB Dataset This dataset contains 42 classes of German traffic signs and we will use Transfer Learning to improve on a vanilla model.

```

[ ]: !wget https://sid.erda.dk/public/archives/daaeac0d7ce1152aea9b61d9f1e19370/
      ↪GTSRB_Final_Training_Images.zip

```

```

--2022-01-27 11:24:48-- https://sid.erda.dk/public/archives/daaeac0d7ce1152aea9
b61d9f1e19370/GTSRB_Final_Training_Images.zip
Resolving sid.erda.dk (sid.erda.dk)... 130.225.104.13
Connecting to sid.erda.dk (sid.erda.dk)|130.225.104.13|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 276294756 (263M) [application/zip]
Saving to: 'GTSRB_Final_Training_Images.zip'

```

```

GTSRB_Final_Trainin 100%[=====>] 263.50M 84.2MB/s in 3.3s

```

```

2022-01-27 11:24:51 (80.5 MB/s) - 'GTSRB_Final_Training_Images.zip' saved
[276294756/276294756]

```

```
[ ]: # Unzipping
!unzip -qq "/GTSRB_Final_Training/Images.zip"
```

```
[ ]: # setting the path to the image folders
IMG_FOLDER = '/content/GTSRB/Final_Training/Images'

# getting all the classes
train_folders = sorted(glob(IMG_FOLDER+'/*'))
# number of classes
print(f"We have {len(train_folders)} classes")
```

We have 43 classes

Helper Functions

```
[ ]: # Helper functions - 3 total !

# 1 - loading the image
def load_image(PATH, resize=True, size=(64,64)):
    """
    Helper function to load the image and resize it to 64x64 in case that is
    needed, or one can also
    define the specific shape
    """
    # loading the image and loading it into RGB
    img = cv2.cvtColor(cv2.imread(PATH), cv2.COLOR_BGR2RGB)

    # resizing and interpolate
    if resize:
        img = cv2.resize(img, (size[0], size[1]), interpolation=cv2.INTER_AREA)

    return img

# 2 - show image
def show_img(PATH):
    """
    Helper function to visualize an image
    """
    # loading the image with the previous function
    img = load_image(PATH)
    # showing the image
    plt.imshow(img)
    plt.axis('off')

# 3 - make grid with different signs
```

```

def make_traffic_grid(PATHS):
    '''
    Helper function which makes a grid to visualize 11 * 4 images (44 total)
    '''
    # loop through the paths we give it and load the images + make into array
    loaded_images = np.array([load_image(img) for img in PATHS])
    # convert it to tensors
    loaded_images = torch.as_tensor(loaded_images)
    # permuting them - changing the RGB --> BGR
    loaded_images = loaded_images.permute(0,3,1,2)
    # make a grid
    img_grid = torchvision.utils.make_grid(loaded_images, nrow=11)
    # make the figure for the grid
    plt.figure(figsize=(24,12))
    # show the image
    plt.imshow(img_grid.permute(1,2,0))
    # turning off the axis
    plt.axis('off')

def make_folders_data(CLASS, DATA_DIR, DATA_SETS):
    '''
    Helper function which makes a folder for the specific classes and the
    specific phase (train, test, val )

    '''
    # data directory
    ddir = Path(DATA_DIR)

    # looping
    for k in DATA_SETS:
        for y in CLASS:
            (ddir/k/y).mkdir(parents=True, exist_ok=True)

def split_data(CLASS_INDECES, CLASSES, DATA_SET, TRAIN_RATIO, VAL_RATIO,
    →DEBUG=True, FILE_BUG=False):
    '''
    Helper function to split the data into 10% testing and validation and 80%
    →training

    '''
    # looping through the index of the classes we chose
    for idx, class_idx in enumerate(CLASS_INDECES):
        # making the image path an array
        img_path = np.array(glob(f"{train_folders[class_idx]}/*.ppm"))
        # getting the specified class name

```

```

class_name = CLASSES[idx]
# debugging
if DEBUG:
    print(f"\nThe class {class_name} has {len(img_path)} images.")
# shuffling them
np.random.shuffle(img_path)
# defining the ratios given the image path length (number of images)
RATIOS = (int(TRAIN_RATIO*len(img_path)), int(VAL_RATIO*len(img_path)))
# splitting them
data_split = np.split(
    img_path,
    indices_or_sections=[RATIOS[0], RATIOS[1]]
)

# zipping
dataset_data = zip(DATA_SET,data_split)
#ddir
data_directory = Path('data')
# moving the specified images
for d, imgs in dataset_data:
    for paths in imgs:
        shutil.copy(paths, f"{data_directory}/{d}/{class_name}")
    if FILE_BUG:
        print(f"Moved {paths} successfully !!!")


def show_image(INPUT,title=None):
    """
    Helper function to visualize the transformations that were done

    """
    # transposing the image
    input = INPUT.numpy().transpose((1,2,0))
    # reverting the normalization
    means = np.array([meanies])
    st_dev = np.array([stds])
    # de-normalize
    input = st_dev * input + means
    # clipping
    input = np.clip(input, 0, 1)
    # showing
    plt.imshow(input)
    if title is not None:
        plt.title(title)
    plt.axis('off')

```

Visualize the Dataset

```
[ ]: # lets visualize a single example
CLASS = 9
sample_image = glob(train_folders[CLASS] + '/*ppm')[0]
show_img(sample_image)
```



```
[ ]: # get the random samples
sample_images = [np.random.choice(glob(f"{x}/*ppm")) for x in train_folders]
print(f"Number of samples : {len(sample_images)}")

# now let's make a grid for these
make_traffic_grid(sample_images)
```

Number of samples : 43



Sepparating the Images

```
[ ]: # we will use specific classes for this
      classes = ['priority_road', 'give_way', 'stop', 'no_entry']

      # and their corresponding indices
      class_index = [12,13,14,17]

      # datasets - train test val
      DATASETS = ['train', 'test', 'val']

      # making folders
      make_folders_data(classes, 'data', DATASETS)

      # splitting the data
      split_data(class_index, classes, DATASETS, 0.8, 0.9, DEBUG=True, FILE_BUG=False)
```

The class priority_road has 2100 images.

The class give_way has 2160 images.

The class stop has 780 images.

The class no_entry has 1110 images.

Image Transformations

```
[ ]: # we will use the standard deviations and means that are given by torchvision
      # means
      meanies = [0.485, 0.456, 0.406]
      # standard deviations
      stds = [0.229, 0.224, 0.225]
```

```

# making a transformations dictionary
# transformations included:
# - Random Resized crop
# - Random Rotation
# - Center Crop
# - Normalization
transforms = {
    "train":T.Compose([
        T.RandomResizedCrop(size=256),
        T.RandomRotation(degrees=15),
        T.RandomHorizontalFlip(),
        T.ToTensor(),
        T.Normalize(meanies, stds)
    ]),

    "val":T.Compose([
        T.Resize(size=256),
        T.CenterCrop(size=224),
        T.ToTensor(),
        T.Normalize(meanies, stds)
    ]),

    "test":T.Compose([
        T.Resize(size=256),
        T.CenterCrop(size=224),
        T.ToTensor(),
        T.Normalize(meanies, stds)
    ])
}

```

Making the datasets

```

[ ]: data_dir = Path('data')
# create a pytorch dataset for each of the images
image_datasets = {
    k: ImageFolder(f"{data_dir}/{k}",transforms[k]) for k in DATASETS
}
# make a dataloader out of the datasets
data_loaders = {
    o: DataLoader(image_datasets[o],batch_size=4,
    shuffle=True,num_workers=2) for o in DATASETS
}

# we can extract the size of the dataset
size_ds = {d: len(image_datasets[d]) for d in DATASETS}


```



```
# also the class names
class_names = image_datasets['train'].classes
```

Visualizing the Image Transformations

```
[ ]: # iterate over one of the dataloaders
targets, labels = next(iter(data_loaders['train']))
# make a grid
out_grid = torchvision.utils.make_grid(targets)
# lets show the images
show_image(out_grid, title=[class_names[x] for x in labels])
```



Making the model - Transfer Learning

```
[ ]: # make model function
def make_model(n_classes):
    '''
    Helper function to use a pretrained model with the GTSRB
    '''
    model = models.resnet34(pretrained=True)
    # stating the number of features
    n_features = model.fc.in_features
    # changing the model fully connected
    model.fc = nn.Linear(n_features, n_classes)
    return model.to(device)

# number of classes
NUM_CLASSES = len(class_names)
# making the base model
base_model = make_model(NUM_CLASSES)
```

Downloading: "<https://download.pytorch.org/models/resnet34-b627a593.pth>" to
/root/.cache/torch/hub/checkpoints/resnet34-b627a593.pth

0%| | 0.00/83.3M [00:00<?, ?B/s]

Training - Testing - Validation Helper Functions

TRAINING

```
[ ]: def train_loop(MODEL, DATA_LOADER, LOSS_FUNCTION, OPTIMIZER, DEVICE, SCHEDULER,
    →EXAMPLES):
    '''
    Helper function which trains the model with the given loss function,
    →optimizer, and scheduler.
    Training on the specified DATA_LOADER and with number of examples
    '''

    # putting into training mode
    model = MODEL.train()
    # accumulate the losses
    running_loss = []
    # correctly predicted counter
    correct_preds = 0
    # going over out dataloader
    for inputs, targets in DATA_LOADER:
        # sending to device
        inputs = inputs.to(DEVICE)
        targets = targets.to(DEVICE)
        # predicting an output
        y_pred = MODEL(inputs)
        # getting the max
        _, preds = torch.max(y_pred, dim=1)
        # getting the loss
        loss = LOSS_FUNCTION(y_pred, targets)
        # correct prediction
        correct_preds+= torch.sum(preds == targets)
        # appending the loss
        running_loss.append(loss.item())
        # zeroing grads
        OPTIMIZER.zero_grad()
        # sending the loss backwards
        loss.backward()
        # stepping
        OPTIMIZER.step()
        #print(num_corr, avg_loss)
    # step the scheduler
    SCHEDULER.step()
    # number correct
    num_corr = correct_preds.double()/EXAMPLES
    avg_loss = np.mean(running_loss)

    return num_corr, avg_loss
```

EVALUTATION

```
[ ]: # Evaluate the model
def val_loop(MODEL, DATA_LOADER, LOSS_FUNCTION, DEVICE, EXAMPLES):
    '''
    Helper function to evaluate the model

    '''
    # setting into eval mode
    model = MODEL.eval()
    # accumulate the loss and predictions
    running_loss = []
    corr_preds = 0
    # loop
    with torch.no_grad():
        for inputs, targets in DATA_LOADER:
            # sending to device
            inputs = inputs.to(DEVICE)
            targets = targets.to(DEVICE)
            # prediction
            y_pred = MODEL(inputs)
            # max vls
            _, preds = torch.max(y_pred, dim=1)
            #loss
            loss = LOSS_FUNCTION(y_pred, targets)
            # correct predictions
            corr_preds += torch.sum(preds == targets)
            # keep track of loss
            running_loss.append(loss.item())
            # correct
            num_corr = corr_preds.double() / EXAMPLES
            avg_loss = np.mean(running_loss)
    return num_corr, avg_loss
```

TRAINING LOOP

```
[ ]: def train_model(MODEL, DATA_LOADER, DATASET_SIZES, DEVICE, EPOCHS=3):
    '''
    Helper Function to train the model with the training and evaluation loop

    '''
    # defin our optimizer
    optimizer = optim.SGD(MODEL.parameters(), lr = 0.001, momentum=0.9)
    # define our scheduler
    scheduler = lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)
    # define our loss function
    loss_fn = nn.CrossEntropyLoss().to(DEVICE)
    # history
    history = defaultdict(list)
```

```

best_acc = 0
for epoch in range(EPOCHS):
    # keep track of epochs
    print(f"\nEpoch: {epoch+1}/{EPOCHS}")
    # TRAINING
    train_acc, train_loss = train_loop(MODEL, DATA_LOADER['train'], loss_fn,
    ↪optimizer, DEVICE, scheduler, DATASET_SIZES['train'])
    train_txt = f"\nTRAINING: LOSS:{train_loss} ACCURACY: {train_acc}"
    print(colored(train_txt, 'red', 'on_white'))
    # VALIDATION
    val_acc, val_loss = val_loop(MODEL, DATA_LOADER['val'], loss_fn, DEVICE,
    ↪DATASET_SIZES['val'])
    val_txt = f"\nVALIDATION: LOSS:{val_loss} ACCURACY: {val_acc}"
    print(colored(val_txt, 'green', 'on_grey'))
    # keeping track of the data
    history['train_acc'].append(train_acc)
    history['train_loss'].append(train_loss)
    history['val_acc'].append(val_acc)
    history['val_loss'].append(val_loss)
    # keeping the best accuracy so far
    if val_acc > best_acc:
        torch.save(MODEL.state_dict(), "best_model_state.bin")
        best_acc = val_acc
    print(f"\nTHE BEST ACCURACY: {best_acc}")
    # loading the best model
    MODEL.load_state_dict(torch.load('best_model_state.bin'))
    return MODEL, history

```

```
[ ]: base_model, history = train_model(base_model, data_loaders, size_ds, device)
```

Epoch: 1/3

```

TRAINING: LOSS:0.3188715517628802 ACCURACY: 0.8876016260162601
VALIDATION: LOSS:0.008109279425806054 ACCURACY: 0.9983739837398374

```

Epoch: 2/3

```

TRAINING: LOSS:0.12497505745169639 ACCURACY: 0.958739837398374
VALIDATION: LOSS:0.0009808485402312986 ACCURACY: 1.0

```

Epoch: 3/3

```

TRAINING: LOSS:0.08614138082549784 ACCURACY: 0.9717479674796748
VALIDATION: LOSS:0.002622378306607988 ACCURACY: 0.9983739837398374


```

THE BEST ACCURACY: 1.0

Visualize the History

```
[ ]: def plot_history(HISTORY,SIZE=(18,6)):
    # creating the figure
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=SIZE)
    # LOSS
    ax1.plot(history['train_loss'], label='Train Loss')
    ax1.plot(history['val_loss'], label='Val Loss')
    # setting the ticks
    ax1.xaxis.set_major_locator(MaxNLocator(integer=True))
    ax1.set_ylim([-0.05, 1.05])
    ax1.legend()
    # labels
    ax1.set_ylabel('Loss')
    ax1.set_xlabel('Epoch')
    # ACCURACY
    ax2.plot(history['train_acc'], label='Train Accuracy')
    ax2.plot(history['val_acc'], label='Val Accuracy')
    # setting the ticks
    ax2.xaxis.set_major_locator(MaxNLocator(integer=True))
    ax2.set_ylim([-0.05, 1.05])
    ax2.legend()
    # labels
    ax2.set_ylabel('Accuracy')
    ax2.set_xlabel('Epoch')
    #Title
    fig.suptitle('Training history')
```

```
[ ]: plot_history(history)
```



Visualize Predictions

```
[ ]: # we are going to show predictions on the test set now
def show_predictions(MODEL, CLASS_NAMES, N_IMAGES=6):
```

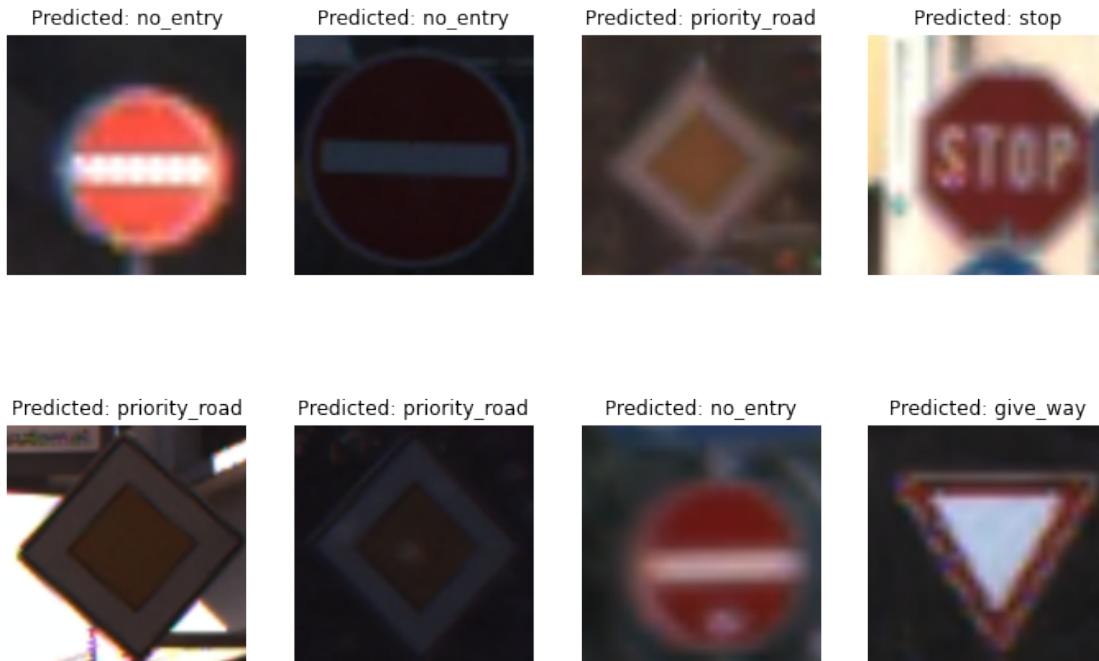
```

# eval mode
MODEL.eval()
# counter
imgs = 0
# create a figure
plt.figure()

# with no grad
with torch.no_grad():
    for idx, (inputs, targets) in enumerate(data_loaders['test']):
        # to device
        inputs = inputs.to(device)
        targets = targets.to(device)
        # preds
        outputs = MODEL(inputs)
        # getting the max
        _, pred = torch.max(outputs, dim=1)
        # now let's check the predictions
        for k in range(inputs.shape[0]):
            imgs+=1
            ax = plt.subplot(2, N_IMAGES//2, imgs)
            ax.set_title(f"Predicted: {class_names[pred[k]]}")
            show_image(inputs.cpu().data[k])
            ax.axis('off')
            if imgs == N_IMAGES:
                return

```

```
[ ]: show_predictions(base_model, class_names, N_IMAGES=8)
```



Getting the actual predictions

```
[ ]: def get_preds(MODEL, DATA_LOADER):
    model = MODEL.eval()
    predictions = []
    real_vals = []
    with torch.no_grad():
        for inputs, targets in DATA_LOADER:
            inputs = inputs.to(device)
            targets = targets.to(device)
            outputs=model(inputs)
            _, preds = torch.max(outputs, dim=1)
            predictions.extend(preds)
            real_vals.extend(targets)
    predictions = torch.as_tensor(predictions).cpu()
    real_vals = torch.as_tensor(real_vals).cpu()
    return predictions, real_vals
```

```
[ ]: y_pred, y_test = get_preds(base_model, data_loaders['test'])
```


Classification Report + Matrix

```
[ ]: # classification report
cr = classification_report(y_test, y_pred, target_names=class_names)
print(cr)
```

	precision	recall	f1-score	support
give_way	1.00	1.00	1.00	216
no_entry	1.00	1.00	1.00	111
priority_road	1.00	1.00	1.00	210
stop	1.00	1.00	1.00	78
accuracy			1.00	615
macro avg	1.00	1.00	1.00	615
weighted avg	1.00	1.00	1.00	615

```
[ ]: # classification matrix
def show_cm(CM, NAMES):
    # copying the confusion matrix
    cm = CM.copy()
    # counts
    cell_counts = cm.flatten()
    # normalize
    cm_row_norm = cm / cm.sum(axis=1)[:, np.newaxis]
    #percentages
    row_percentages = ["{0:.2f}".format(value) for value in cm_row_norm.flatten()]
    # labels
    cell_labels = [f"{cnt}\n{per}" for cnt, per in zip(cell_counts,
    ↪row_percentages)]
    cell_labels = np.asarray(cell_labels).reshape(cm.shape[0], cm.shape[1])
    # dataframe
    df = pd.DataFrame(cm_row_norm, index=NAMES, columns=NAMES)
    #heatmap
    hmap = sns.heatmap(df, annot=cell_labels, fmt="", cmap="Blues")
    # ticks
    hmap.yaxis.set_ticklabels(hmap.yaxis.get_ticklabels(), rotation=0, ha='right')
    hmap.xaxis.set_ticklabels(hmap.xaxis.get_ticklabels(), rotation=30,
    ↪ha='right')
    # labels
    plt.ylabel('True Sign')
    plt.xlabel('Predicted Sign')
```

```
[ ]: # make the confusion matrix
cm = confusion_matrix(y_test, y_pred)
# showing it
show_cm(cm, class_names)
```

Predict on Unseen Images

```
[ ]: # we need to get the probabilities for each of the classes
```

```
def get_pred_proba(MODEL, PATH):
    # reading and converting the images
    img = Image.open(PATH)
    img = img.convert('RGB')
    img = transforms['test'](img).unsqueeze(0)
    img = img.to(device)
    # predicting
    pred = MODEL(img)
    pred = F.softmax(pred, dim=1)
    pred = pred.detach().cpu().numpy().flatten()
    return pred
```

```
# define another function to plot it as it returns an array
```

```
def plot_pred_proba(PRED, NAMES):
    pred_df = pd.DataFrame(
        {
            'Class_Name': NAMES,
            'Values': PRED
        })
```

```
sns.barplot(x='Values',y='Class_Name',data=pred_df,orient='h')
plt.xlim([0,1])
plt.show()
```

```
[ ]: !gdown --id 19Qz3a610u_QSHsLeTznx8LtDBu4tbqHr
!gdown --id 1F61-iNh1Jk-yKZRGcu6S9P29HxDFxF0u
```

Downloading...

From: https://drive.google.com/uc?id=19Qz3a610u_QSHsLeTznx8LtDBu4tbqHr

To: /content/stop-sign.jpg

100% 77.3k/77.3k [00:00<00:00, 28.8MB/s]


Downloading...

From: <https://drive.google.com/uc?id=1F61-iNh1Jk-yKZRGcu6S9P29HxDFxF0u>


To: /content/unknown-sign.jpg

100% 41.4k/41.4k [00:00<00:00, 38.8MB/s]

```
[ ]: # getting the prediction and then visualizing it
predz = get_pred_proba(base_model, 'stop-sign.jpg')
# visualize
plot_pred_proba(predz, class_names)
```



```
[ ]: # getting the prediction and then visualizing it
predz = get_pred_proba(base_model, 'unknown-sign.jpg')
# visualize
plot_pred_proba(predz, class_names)
```



Adding the Unknown class Dan Hendrycks and Kevin Gimpel demonstrate that we can measure whether a sample is misclassified by looking at the in-sample and out-of sample distributions of the softmax classification. This can cause that Neural Networks confidently predict a wrong class when the sample is not part of the training distribution. We see that different training and testing distributions affect the quality of the neural network. The paper can be found here:

<https://arxiv.org/pdf/1610.02136.pdf>

```
[ ]: # First get the indices of the examples that we did not use
not_used_index = [x for x, _ in enumerate(train_folders) if x not in
    ↪class_index]
print(len(not_used_index))
```

39

```
[ ]: def make_folders_move_unknowns():
    '''
    Helper to move the unknowns and make a folder

    '''
    # make the folder for the datasets with the label unknown
    for sets in DATASETS:
        (data_dir/sets/'unknown').mkdir(parents=True, exist_ok=True)

    # now we find the images
```

```

for uking in not_used_index:
    img_path = np.array(glob(f"{train_folders[uking]}/*.ppm"))
    img_path = np.random.choice(img_path, 200)
    # splitting the data
    train_split = int(0.8*len(img_path)) # keeping 80% for training
    test_split = int(0.9*len(img_path)) # keeping 10 % for val + 10% testing
    ds_split = np.split(img_path, indices_or_sections=[train_split, test_split])
    # zipping
    dataset_dt = zip(DATASETS, ds_split)
    # moving the images
    for ds, imgs in dataset_dt:
        for paths in imgs:
            shutil.copy(paths, f"{data_dir}/{ds}/unknown/")

make_folders_move_unknowns()

```

```

[ ]: # transform the dataset
x_image_dataset = {
    d: ImageFolder(f"{data_dir}/{d}", transforms[d]) for d in DATASETS
}

# dataloader
data_loaders = {

    d: DataLoader(x_image_dataset[d], batch_size=4, num_workers=2,
    ↪shuffle=True) for d in DATASETS
}

# getting the sizes of the datasets
dataset_sizes = {
    d: len(x_image_dataset[d]) for d in DATASETS
}

# and finally the class names
x_class_names = x_image_dataset['train'].classes

# lets look at the sizes now
dataset_sizes

```

```

[ ]: {'test': 648, 'train': 5086, 'val': 646}

```

Test new model with extra class

```

[ ]: # new model with new features
x_n_features = len(x_class_names)
x_model = make_model(x_n_features)

```

```
# training the new model
```

```
x_model, x_history = train_model(x_model, data_loaders, dataset_sizes, device)
```

Epoch: 1/3

```
TRAINING: LOSS:0.35815324049461583 ACCURACY: 0.8786865906409752
```

```
VALIDATION: LOSS:0.010733573848427361 ACCURACY: 0.9969040247678019
```

Epoch: 2/3

```
TRAINING: LOSS:0.139202132547594 ACCURACY: 0.951042076287849
```

```
VALIDATION: LOSS:0.005631997385669955 ACCURACY: 0.998452012383901
```

Epoch: 3/3

```
TRAINING: LOSS:0.08849910689006274 ACCURACY: 0.9732599292174596
```


```
VALIDATION: LOSS:0.014248965154372947 ACCURACY: 0.9969040247678019
```

THE BEST ACCURACY: 0.998452012383901

```
[ ]: # we can plot the new training history  
plot_history(x_history)
```



```
[ ]: # getting the prediction and then visualizing it  
x_predz = get_pred_proba(x_model, 'unknown-sign.jpg')  
# visualize  
plot_pred_proba(x_predz, x_class_names)
```



```
[ ]: # lets look at the examples in our dataset
show_predictions(x_model, x_class_names, 8)
```



```
[ ]: # we can make the classification again
# get the predictions
y_pred, y_test = get_preds(x_model, data_loaders['test'])

# make the classification report
# classification report
cr = classification_report(y_test, y_pred, target_names=x_class_names)
print(cr)

# show the confusion matrix
# make the confusion matrix
cm = confusion_matrix(y_test, y_pred)
# showing it
show_cm(cm, x_class_names)
```

	precision	recall	f1-score	support
give_way	1.00	1.00	1.00	216
no_entry	1.00	1.00	1.00	111
priority_road	1.00	1.00	1.00	210
stop	1.00	1.00	1.00	78
unknown	1.00	1.00	1.00	33
accuracy			1.00	648
macro avg	1.00	1.00	1.00	648
weighted avg	1.00	1.00	1.00	648

