

Chapter_3_TL

January 26, 2022

1 Chapter 3: Transfer Learning - German Traffic Sign Recognition Benchmark (GTSRB)

Imports

```
[ ]: #files
from glob import glob
import shutil
from pathlib import Path

# linalg
import numpy as np

# images
import cv2

# plotting
from matplotlib.ticker import MaxNLocator
import matplotlib.pyplot as plt
import seaborn as sns
from pylab import rcParams

# dataframes
import pandas as pd

# sklearn
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report

#progress bar
from tqdm import tqdm

# nice colors
import termcolor
from termcolor import colored

# pytorch
import torch, torchvision
```

```

from torch import nn, optim

# functional
import torch.nn.functional as F

# utils
import torchvision.transforms as T
import PIL.Image as Image
from torch.optim import lr_scheduler
from collections import defaultdict

# datasets
from torchvision.datasets import ImageFolder
from torch.utils.data import DataLoader

# models
from torchvision import models

```

```

[ ]: # getting the specific device
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
dev_name = torch.cuda.get_device_name()
print(dev_name)

```

Tesla K80

```

[ ]: # setting the figure size
rcParams['figure.figsize'] = 12,8

```

Getting the GTSRB Dataset This dataset contains 42 classes of German traffic signs and we will use Transfer Learning to improve on a vanilla model.

```

[ ]: !wget https://sid.erda.dk/public/archives/daaeac0d7ce1152aea9b61d9f1e19370/
      ↪GTSRB_Final_Training_Images.zip

```

```

--2022-01-26 18:26:49-- https://sid.erda.dk/public/archives/daaeac0d7ce1152aea9
b61d9f1e19370/GTSRB_Final_Training_Images.zip
Resolving sid.erda.dk (sid.erda.dk)... 130.225.104.13
Connecting to sid.erda.dk (sid.erda.dk)|130.225.104.13|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 276294756 (263M) [application/zip]
Saving to: 'GTSRB_Final_Training_Images.zip'

```

```

GTSRB_Final_Trainin 100%[=====>] 263.50M  11.1MB/s    in 26s

```

```

2022-01-26 18:27:16 (9.99 MB/s) - 'GTSRB_Final_Training_Images.zip' saved
[276294756/276294756]

```

```
[ ]: # Unzipping
!unzip -qq "/GTSRB_Final_Training/Images.zip"
```

```
[ ]: # setting the path to the image folders
IMG_FOLDER = '/content/GTSRB/Final_Training/Images'

# getting all the classes
train_folders = sorted(glob(IMG_FOLDER+'/*'))
# number of classes
print(f"We have {len(train_folders)} classes")
```

We have 43 classes

Helper Functions

```
[ ]: # Helper functions - 3 total !

# 1 - loading the image
def load_image(PATH, resize=True, size=(64,64)):
    """
    Helper function to load the image and resize it to 64x64 in case that is
    needed, or one can also
    define the specific shape
    """
    # loading the image and loading it into RGB
    img = cv2.cvtColor(cv2.imread(PATH), cv2.COLOR_BGR2RGB)

    # resizing and interpolate
    if resize:
        img = cv2.resize(img, (size[0], size[1]), interpolation=cv2.INTER_AREA)

    return img

# 2 - show image
def show_img(PATH):
    """
    Helper function to visualize an image
    """
    # loading the image with the previous function
    img = load_image(PATH)
    # showing the image
    plt.imshow(img)
    plt.axis('off')

# 3 - make grid with different signs
```

```

def make_traffic_grid(PATHS):
    '''
    Helper function which makes a grid to visualize 11 * 4 images (44 total)
    '''
    # loop through the paths we give it and load the images + make into array
    loaded_images = np.array([load_image(img) for img in PATHS])
    # convert it to tensors
    loaded_images = torch.as_tensor(loaded_images)
    # permuting them - changing the RGB --> BGR
    loaded_images = loaded_images.permute(0,3,1,2)
    # make a grid
    img_grid = torchvision.utils.make_grid(loaded_images, nrow=11)
    # make the figure for the grid
    plt.figure(figsize=(24,12))
    # show the image
    plt.imshow(img_grid.permute(1,2,0))
    # turning off the axis
    plt.axis('off')

def make_folders_data(CLASS, DATA_DIR, DATA_SETS):
    '''
    Helper function which makes a folder for the specific classes and the
    specific phase (train, test, val )

    '''
    # data directory
    ddir = Path(DATA_DIR)

    # looping
    for k in DATA_SETS:
        for y in CLASS:
            (ddir/k/y).mkdir(parents=True, exist_ok=True)

def split_data(CLASS_INDECES, CLASSES, DATA_SET, TRAIN_RATIO, VAL_RATIO,
    →DEBUG=True, FILE_BUG=False):
    '''
    Helper function to split the data into 10% testing and validation and 80%
    →training

    '''

    # looping through the index of the classes we chose
    for idx, class_idx in enumerate(CLASS_INDECES):
        # making the image path an array
        img_path = np.array(glob(f"{train_folders[class_idx]}/*.ppm"))
        # getting the specified class name

```

```

class_name = CLASSES[idx]
# debugging
if DEBUG:
    print(f"\nThe class {class_name} has {len(img_path)} images.")
# shuffling them
np.random.shuffle(img_path)
# defining the ratios given the image path length (number of images)
RATIOS = (int(TRAIN_RATIO*len(img_path)), int(VAL_RATIO*len(img_path)))
# splitting them
data_split = np.split(
    img_path,
    indices_or_sections=[RATIOS[0], RATIOS[1]]
)

# zipping
dataset_data = zip(DATA_SET,data_split)
#ddir
data_directory = Path('data')
# moving the specified images
for d, imgs in dataset_data:
    for paths in imgs:
        shutil.copy(paths, f"{data_directory}/{d}/{class_name}")
    if FILE_BUG:
        print(f"Moved {paths} successfully !!!")


def show_image(INPUT,title=None):
    """
    Helper function to visualize the transformations that were done

    """
    # transposing the image
    input = INPUT.numpy().transpose((1,2,0))
    # reverting the normalization
    means = np.array([meanies])
    st_dev = np.array([stds])
    # de-normalize
    input = st_dev * input + means
    # clipping
    input = np.clip(input, 0, 1)
    # showing
    plt.imshow(input)
    if title is not None:
        plt.title(title)
    plt.axis('off')

```

Visualize the Dataset

```
[ ]: # lets visualize a single example
CLASS = 9
sample_image = glob(train_folders[CLASS] + '/*ppm')[0]
show_img(sample_image)
```



```
[ ]: # get the random samples
sample_images = [np.random.choice(glob(f"{x}/*ppm")) for x in train_folders]
print(f"Number of samples : {len(sample_images)}")

# now let's make a grid for these
make_traffic_grid(sample_images)
```

Number of samples : 43



Sepparating the Images

```
[ ]: # we will use specific classes for this
classes = ['priority_road', 'give_way', 'stop', 'no_entry']

# and their corresponding indices
class_index = [12,13,14,17]

# datasets - train test val
DATASETS = ['train', 'test', 'val']

# making folders
make_folders_data(classes, 'data', DATASETS)

# splitting the data
split_data(class_index, classes, DATASETS, 0.8, 0.9, DEBUG=True, FILE_BUG=False)
```

The class priority_road has 2100 images.

The class give_way has 2160 images.

The class stop has 780 images.

The class no_entry has 1110 images.

Image Transformations

```
[ ]: # we will use the standard deviations and means that are given by torchvision
# means
meanies = [0.485, 0.456, 0.406]
# standard deviations
stds = [0.229, 0.224, 0.225]
```

```

# making a transformations dictionary
# transformations included:
# - Random Resized crop
# - Random Rotation
# - Center Crop
# - Normalization
transforms = {
    "train": T.Compose([
        T.RandomResizedCrop(size=256),
        T.RandomRotation(degrees=15),
        T.RandomHorizontalFlip(),
        T.ToTensor(),
        T.Normalize(meanies, stds)
    ]),

    "val": T.Compose([
        T.Resize(size=256),
        T.CenterCrop(size=224),
        T.ToTensor(),
        T.Normalize(meanies, stds)
    ]),

    "test": T.Compose([
        T.Resize(size=256),
        T.CenterCrop(size=224),
        T.ToTensor(),
        T.Normalize(meanies, stds)
    ])
}

```

Making the datasets

```

[ ]: data_dir = Path('data')
# create a pytorch dataset for each of the images
image_datasets = {
    k: ImageFolder(f"{data_dir}/{k}", transforms[k]) for k in DATASETS
}
# make a dataloader out of the datasets
data_loaders = {
    o: DataLoader(image_datasets[o], batch_size=4,
        shuffle=True, num_workers=2) for o in DATASETS
}

# we can extract the size of the dataset
size_ds = {d: len(image_datasets[d]) for d in DATASETS}

```



```
# also the class names
class_names = image_datasets['train'].classes
```

Visualizing the Image Transformations

```
[ ]: # iterate over one of the dataloaders
targets, labels = next(iter(data_loaders['train']))
# make a grid
out_grid = torchvision.utils.make_grid(targets)
# lets show the images
show_image(out_grid, title=[class_names[x] for x in labels])
```



Making the model - Transfer Learning

```
[ ]: # make model function
def make_model(n_classes):
    '''
    Helper function to use a pretrained model with the GTSRB
    '''
    model = models.resnet34(pretrained=True)
    # stating the number of features
    n_features = model.fc.in_features
    # changing the model fully connected
    model.fc = nn.Linear(n_features, n_classes)
    return model.to(device)

# number of classes
NUM_CLASSES = len(class_names)
# making the base model
base_model = make_model(NUM_CLASSES)
```

Downloading: "<https://download.pytorch.org/models/resnet34-b627a593.pth>" to
/root/.cache/torch/hub/checkpoints/resnet34-b627a593.pth

0%| | 0.00/83.3M [00:00<?, ?B/s]

Training - Testing - Validation Helper Functions

2 TRAINING

```
[ ]: def train_loop(MODEL, DATA_LOADER, LOSS_FUNCTION, OPTIMIZER, DEVICE, SCHEDULER,
    ↪EXAMPLES):
    '''
    Helper function which trains the model with the given loss function,
    ↪optimizer, and scheduler.
    Training on the specified DATA_LOADER and with number of examples
    '''

    # putting into training mode
    model = MODEL.train()
    # accumulate the losses
    running_loss = []
    # correctly predicted counter
    correct_preds = 0
    # going over out dataloader
    for inputs, targets in DATA_LOADER:
        # sending to device
        inputs = inputs.to(DEVICE)
        targets = targets.to(DEVICE)
        # predicting an output
        y_pred = MODEL(inputs)
        # getting the max
        _, preds = torch.max(y_pred, dim=1)
        # getting the loss
        loss = LOSS_FUNCTION(y_pred, targets)
        # correct prediction
        correct_preds+= torch.sum(preds == targets)
        # appending the loss
        running_loss.append(loss.item())
        # zeroing grads
        OPTIMIZER.zero_grad()
        # sending the loss backwards
        loss.backward()
        # stepping
        OPTIMIZER.step()
        #print(num_corr, avg_loss)
    # step the scheduler
    SCHEDULER.step()
    # number correct
    num_corr = correct_preds.double()/EXAMPLES
    avg_loss = np.mean(running_loss)

    return num_corr, avg_loss
```

3 EVALUTATION

```
[ ]: # Evaluate the model
def val_loop(MODEL, DATA_LOADER, LOSS_FUNCTION, DEVICE, EXAMPLES):
    '''
    Helper function to evaluate the model

    '''
    # setting into eval mode
    model = MODEL.eval()
    # accumulate the loss and predictions
    running_loss = []
    corr_preds = 0
    # loop
    with torch.no_grad():
        for inputs, targets in DATA_LOADER:
            # sending to device
            inputs = inputs.to(DEVICE)
            targets = targets.to(DEVICE)
            # prediction
            y_pred = MODEL(inputs)
            # max vls
            _, preds = torch.max(y_pred, dim=1)
            #loss
            loss = LOSS_FUNCTION(y_pred, targets)
            # correct predictions
            corr_preds += torch.sum(preds == targets)
            # keep track of loss
            running_loss.append(loss.item())
            # correct
            num_corr = corr_preds.double() / EXAMPLES
            avg_loss = np.mean(running_loss)
    return num_corr, avg_loss
```

4 TRAINING LOOP

```
[ ]: def train_model(MODEL, DATA_LOADER, DATASET_SIZES, DEVICE, EPOCHS=3):
    '''
    Helper Function to train the model with the training and evaluation loop

    '''
    # defin our optimizer
    optimizer = optim.SGD(MODEL.parameters(), lr = 0.001, momentum=0.9)
    # define our scheduler
    scheduler = lr_scheduler.StepLR(optimizer, step_size=7, gamma=0.1)
    # define our loss function
```

```

loss_fn = nn.CrossEntropyLoss().to(DEVICE)
# history
history = defaultdict(list)
best_acc = 0
for epoch in range(EPOCHS):
    # keep track of epochs
    print(f"\nEpoch: {epoch+1}/{EPOCHS}")
    # TRAINING
    train_acc, train_loss = train_loop(MODEL, DATA_LOADER['train'], loss_fn,
    ↪optimizer, DEVICE, scheduler, DATASET_SIZES['train'])
    train_txt = f"\nTRAINING: LOSS:{train_loss} ACCURACY: {train_acc}"
    print(colored(train_txt, 'red', 'on_white'))
    # VALIDATION
    val_acc, val_loss = val_loop(MODEL, DATA_LOADER['val'], loss_fn, DEVICE,
    ↪DATASET_SIZES['val'])
    val_txt = f"\nVALIDATION: LOSS:{val_loss} ACCURACY: {val_acc}"
    print(colored(val_txt, 'green', 'on_grey'))
    # keeping track of the data
    history['train_acc'].append(train_acc)
    history['train_loss'].append(train_loss)
    history['val_acc'].append(val_acc)
    history['val_loss'].append(val_loss)
    # keeping the best accuracy so far
    if val_acc > best_acc:
        torch.save(MODEL.state_dict(), "best_model_state.bin")
        best_acc = val_acc
    print(f"\nTHE BEST ACCURACY: {best_acc}")
    # loading the best model
    MODEL.load_state_dict(torch.load('best_model_state.bin'))
    return MODEL, history

```

```
[ ]: base_model, history = train_model(base_model, data_loaders, size_ds, device)
```

Epoch: 1/3

```

TRAINING: LOSS:0.05681165572400275 ACCURACY: 0.9817073170731707
VALIDATION: LOSS:0.004216150927070821 ACCURACY: 0.9983739837398374

```

Epoch: 2/3

```

TRAINING: LOSS:0.054183467861499905 ACCURACY: 0.983130081300813
VALIDATION: LOSS:0.0028477875562376307 ACCURACY: 1.0

```

Epoch: 3/3

```

TRAINING: LOSS:0.0534301981483763 ACCURACY: 0.9853658536585366
VALIDATION: LOSS:0.0037492613879549423 ACCURACY: 0.9983739837398374

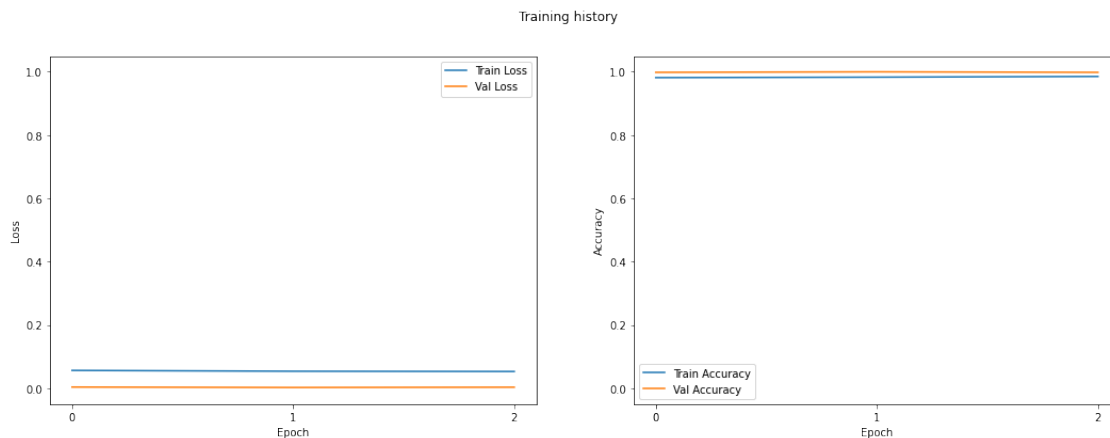
```

THE BEST ACCURACY: 1.0

Visualize the History

```
[ ]: def plot_history(HISTORY,SIZE=(18,6)):  
    # creating the figure  
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=SIZE)  
    # LOSS  
    ax1.plot(history['train_loss'], label='Train Loss')  
    ax1.plot(history['val_loss'], label='Val Loss')  
    # setting the ticks  
    ax1.xaxis.set_major_locator(MaxNLocator(integer=True))  
    ax1.set_ylim([-0.05, 1.05])  
    ax1.legend()  
    # labels  
    ax1.set_ylabel('Loss')  
    ax1.set_xlabel('Epoch')  
    # ACCURACY  
    ax2.plot(history['train_acc'], label='Train Accuracy')  
    ax2.plot(history['val_acc'], label='Val Accuracy')  
    # setting the ticks  
    ax2.xaxis.set_major_locator(MaxNLocator(integer=True))  
    ax2.set_ylim([-0.05, 1.05])  
    ax2.legend()  
    # labels  
    ax2.set_ylabel('Accuracy')  
    ax2.set_xlabel('Epoch')  
    #Title  
    fig.suptitle('Training history')
```

```
[ ]: plot_history(history)
```



Visualize Predictions

```
[ ]: # we are going to show predictions on the test set now
def show_predictions(MODEL, CLASS_NAMES, N_IMAGES=6):
    # eval mode
    MODEL.eval()
    # counter
    imgs = 0
    # create a figure
    plt.figure()

    # with no grad
    with torch.no_grad():
        for idx, (inputs, targets) in enumerate(data_loaders['test']):
            # to device
            inputs = inputs.to(device)
            targets = targets.to(device)
            #preds
            outputs = MODEL(inputs)
            # getting the max
            _, pred = torch.max(outputs, dim=1)
            # now let's check the predictions
            for k in range(inputs.shape[0]):
                imgs+=1
                ax = plt.subplot(2, N_IMAGES//2, imgs)
                ax.set_title(f"Predicted: {class_names[pred[k]]}")
                show_image(inputs.cpu().data[k])
                ax.axis('off')
            if imgs == N_IMAGES:
                return
```

```
[ ]: show_predictions(base_model, class_names, N_IMAGES=8)
```



Getting the actual predictions

```
[ ]: def get_preds(MODEL, DATA_LOADER):
    model = MODEL.eval()
    predictions = []
    real_vals = []
    with torch.no_grad():
        for inputs, targets in DATA_LOADER:
            inputs = inputs.to(device)
            targets = targets.to(device)
            outputs=model(inputs)
            _, preds = torch.max(outputs, dim=1)
            predictions.extend(preds)
            real_vals.extend(targets)
    predictions = torch.as_tensor(predictions).cpu()
    real_vals = torch.as_tensor(real_vals).cpu()
    return predictions, real_vals
```

```
[ ]: y_pred, y_test = get_preds(base_model, data_loaders['test'])
```


Classification Report + Matrix

```
[ ]: # classification report
cr = classification_report(y_test, y_pred, target_names=class_names)
print(cr)
```

	precision	recall	f1-score	support
give_way	1.00	1.00	1.00	216
no_entry	1.00	0.99	1.00	111
priority_road	1.00	1.00	1.00	210
stop	1.00	1.00	1.00	78
accuracy			1.00	615
macro avg	1.00	1.00	1.00	615
weighted avg	1.00	1.00	1.00	615

```
[ ]: # classification matrix
def show_cm(CM, NAMES):
    # copying the confusion matrix
    cm = CM.copy()
    # counts
    cell_counts = cm.flatten()
    # normalize
    cm_row_norm = cm / cm.sum(axis=1)[:, np.newaxis]
    #percentages
    row_percentages = ["{0:.2f}".format(value) for value in cm_row_norm.flatten()]
    # labels
    cell_labels = [f"{cnt}\n{per}" for cnt, per in zip(cell_counts,
    ↪row_percentages)]
    cell_labels = np.asarray(cell_labels).reshape(cm.shape[0], cm.shape[1])
    # dataframe
    df = pd.DataFrame(cm_row_norm, index=NAMES, columns=NAMES)
    #heatmap
    hmap = sns.heatmap(df, annot=cell_labels, fmt="", cmap="Blues")
    # ticks
    hmap.yaxis.set_ticklabels(hmap.yaxis.get_ticklabels(), rotation=0, ha='right')
    hmap.xaxis.set_ticklabels(hmap.xaxis.get_ticklabels(), rotation=30,
    ↪ha='right')
    # labels
    plt.ylabel('True Sign')
    plt.xlabel('Predicted Sign')
```

```
[ ]: # make the confusion matrix
cm = confusion_matrix(y_test, y_pred)
# showing it
show_cm(cm, class_names)
```

5 Predict on Unseen Images

```
[ ]: # we need to get the probabilities for each of the classes
def get_pred_proba(MODEL, PATH):
    # reading and converting the images
    img = Image.open(PATH)
    img = img.convert('RGB')
    img = transforms['test'](img).unsqueeze(0)
    img = img.to(device)
    # predicting
    pred = MODEL(img)
    pred = F.softmax(pred, dim=1)
    pred = pred.detach().cpu().numpy().flatten()
    return pred

# define another function to plot it as it returns an array
def plot_pred_proba(PRED, NAMES):
    pred_df = pd.DataFrame(
        {
            'Class_Name': NAMES,
```

```

        'Values':PRED
    })
sns.barplot(x='Values',y='Class_Name',data=pred_df,orient='h')
plt.xlim([0,1])
plt.show()

```

```

[ ]: !gdown --id 19Qz3a610u_QSHsLeTznx8LtDBu4tbqHr
!gdown --id 1F61-iNh1Jk-yKZRGcu6S9P29HxDFxF0u

```

Downloading...

From: https://drive.google.com/uc?id=19Qz3a610u_QSHsLeTznx8LtDBu4tbqHr

To: /content/stop-sign.jpg

100% 77.3k/77.3k [00:00<00:00, 73.7MB/s]

Downloading...

From: <https://drive.google.com/uc?id=1F61-iNh1Jk-yKZRGcu6S9P29HxDFxF0u>


To: /content/unknown-sign.jpg

100% 41.4k/41.4k [00:00<00:00, 16.1MB/s]


```

[ ]: # getting the prediction and then visualizing it
predz = get_pred_proba(base_model, 'stop-sign.jpg')
# visualize
plot_pred_proba(predz, class_names)

```



```
[ ]: # getting the prediction and then visualizing it
predz = get_pred_proba(base_model, 'unknown-sign.jpg')
# visualize
plot_pred_proba(predz, class_names)
```



```
[ ]:
```