

Chapter_4_LSTM

January 27, 2022

1 Chapter 4: Time Series Forecasting with LSTMs

Imports

```
[ ]: # OS
import os

# plotting
import matplotlib.pyplot as plt
import seaborn
from pylab import rcParams

# linalg
import numpy as np

# dfs
import pandas as pd

# preprocess
from sklearn.preprocessing import MinMaxScaler

# colored print
import termcolor
from termcolor import colored

# torch
import torch
from torch import nn, optim
```

```
[ ]: rcParams['figure.figsize'] = 18,10
```

Fetch the dataset

```
[ ]: !gdown --id 1AsfdLrGESQnRW5rbMz56A1KBc3Fe5aV
```

Downloading...

From: <https://drive.google.com/uc?id=1AsfdLrGESQnRW5rbMz56A1KBc3Fe5aV>

To: /content/time_series_19-covid-Confirmed.csv

100% 19.2k/19.2k [00:00<00:00, 7.15MB/s]

```
[ ]: !wget https://raw.githubusercontent.com/CSSEGISandData/COVID-19/master/  
      ↳csse_covid_19_data/csse_covid_19_time_series/  
      ↳time_series_covid19_confirmed_global.csv
```

```
--2022-01-27 15:25:37-- https://raw.githubusercontent.com/CSSEGISandData/COVID-  
19/master/csse_covid_19_data/csse_covid_19_time_series/time_series_covid19_confirmed_global.csv
```

```
Resolving raw.githubusercontent.com (raw.githubusercontent.com)...
```

```
185.199.108.133, 185.199.109.133, 185.199.110.133, ...
```

```
Connecting to raw.githubusercontent.com
```

```
(raw.githubusercontent.com)|185.199.108.133|:443... failed: Connection timed  
out.
```

```
Connecting to raw.githubusercontent.com
```

```
(raw.githubusercontent.com)|185.199.109.133|:443... connected.
```

```
HTTP request sent, awaiting response... 200 OK
```

```
Length: 1050571 (1.0M) [text/plain]
```

```
Saving to: 'time_series_covid19_confirmed_global.csv.1'
```


```
time_series_covid19 100%[=====>] 1.00M --.-KB/s in 0.06s
```

```
2022-01-27 15:27:47 (17.3 MB/s) - 'time_series_covid19_confirmed_global.csv.1'  
saved [1050571/1050571]
```

Load and view dataset


```
[ ]: # assign path  
PATH_DATA = '/content/time_series_covid19_confirmed_global.csv'  
# reading  
df = pd.read_csv(PATH_DATA)  
  
# we need to get rid of the first four columns:  
# PROVINCE, COUNTRY, LAT & LONG --> Not much information except prediction per_  
↳country ?  
df = df.iloc[:, 4:]  
  
# check the nans == 0  
assert df.isnull().sum().sum() == 0  
  
# lets do a cummulative sum  
daily_cases = df.sum(axis=0)  
# get the index of the daily cases and convert to timestamp  
daily_cases.index = pd.to_datetime(daily_cases.index)  
# lets look at it  
daily_cases.head()  
# lets visualize it  
plt.plot(daily_cases);  
plt.title("Graph 1. Cummulative Daily Cases");
```

```
plt.show()
```



Preprocessing: Removing cumulative

```
[ ]: # removing the cumulative by removing the current value from the previous
daily_cases = daily_cases.diff().fillna(daily_cases[0]).astype(np.int64)
daily_cases.head()
plt.plot(daily_cases);
plt.title("Graph 2. Daily Cases");
plt.show()
```



```
[ ]: # lets reserve 90% for training, 10% for testing
test_size = round(len(daily_cases)*0.1)
#test_size=14

# the training data
train_data = daily_cases[:-test_size]
test_data = daily_cases[-test_size:]

#check the shape
#print(train_data.shape, test_data.shape)

# we also need to scale our data
scaler = MinMaxScaler()
# fitting to the data
scaler = scaler.fit(np.expand_dims(train_data, axis=1))
# quick func
def scale_data(SCALE, DATA): return SCALE.transform(np.expand_dims(DATA,
↪axis=1))

# transforming the previous data
train_data = scale_data(scaler, train_data)
test_data = scale_data(scaler, test_data)

# we also need to chunk the sequences, because they are continious
def create_sequences(DATA, LENGTH):
    xs,ys = [],[]
    for k in range(len(DATA)-LENGTH-1):
```

```

        x = DATA[k:(k+LENGTH)]
        y = DATA[k+LENGTH]
        xs.append(x), ys.append(y)
    return np.array(xs), np.array(ys)

# make the defined sequences
SEQ_LENGTH = 5
# the data we will use
X_train, y_train = create_sequences(train_data, SEQ_LENGTH)
X_test, y_test = create_sequences(test_data, SEQ_LENGTH)
# converting the vals
# TRAINING
X_train = torch.from_numpy(X_train).float()
y_train = torch.from_numpy(y_train).float()
# TESTING
X_test = torch.from_numpy(X_test).float()
y_test = torch.from_numpy(y_test).float()

# lets look at the shape of each of these samples
X_train.shape

```

```
[ ]: torch.Size([656, 5, 1])
```

Build the model !

```

[ ]: # building the LST class
# we pass in the number of features, number of hidden units, the sequence_
    ↳length and the number of layers
class nCOV_LSTM(nn.Module):
    def __init__(self, n_features, n_hidden, seq_len, DEVICE, n_layers=2):
        super(nCOV_LSTM, self).__init__()
        # the values
        self.device = DEVICE
        self.n_hidden = n_hidden
        self.seq_len = seq_len
        self.n_layers = n_layers
        # the actual LSTM
        self.lstm = nn.LSTM(
            input_size=n_features,
            hidden_size = n_hidden,
            num_layers = n_layers,
            dropout=0.5
        )
        # the linear layer: in--> number of hidden units, and out is one bc we are_
    ↳predicting a single value
        self.linear = nn.Linear(in_features=n_hidden, out_features=1)

```

```

# we need to reset the state after each sample --> stateless LSTM
def reset_hidden_state(self):
    self.hidden = (
        torch.zeros(self.n_layers, self.seq_len, self.n_hidden, device=self.
→device),
        torch.zeros(self.n_layers, self.seq_len, self.n_hidden, device=self.
→device)
    )
# defining the forwards pass
def forward(self, sequences):
    lstm_out, self.hidden = self.lstm(
        sequences.view(len(sequences), self.
→seq_len, -1),
        self.hidden)
    last_time_step = lstm_out.view(
        self.seq_len, len(sequences), self.n_hidden
    )[-1]
    y_pred = self.linear(last_time_step)
    return y_pred

```

Training Loop

```

[ ]: def train_LSTM_model(MODEL, TRAIN_DATA,
→TRAIN_LABELS,EPOCHS,DEVICE,TEST_DATA=None, TEST_LABELS=None):
    # sending them all to device?
    TRAIN_DATA = TRAIN_DATA.to(DEVICE)
    TRAIN_LABELS = TRAIN_LABELS.to(DEVICE)
    if (TEST_DATA is not None) and (TEST_LABELS is not None):
        TEST_DATA = TEST_DATA.to(DEVICE)
        TEST_LABELS = TEST_LABELS.to(DEVICE)
    # define our loss function
    loss_func = nn.MSELoss(reduction='sum')
    # optimizer
    opt = optim.Adam(MODEL.parameters(), lr=1e-3)
    # history tracker
    train_history,test_history = np.zeros(EPOCHS), np.zeros(EPOCHS)
    # looping through the epochs
    for epoch in range(EPOCHS):
        # putting into train mode
        #MODEL.train()
        # reset the state just in case
        MODEL.reset_hidden_state()
        # predicting
        y_pred = MODEL(TRAIN_DATA)
        # calculate the loss
        loss = loss_func(y_pred.float(), TRAIN_LABELS)
        # appending the training history

```

```

train_history[epoch] = loss.item()
# testing
if TEST_DATA is not None:
    with torch.no_grad():
        #MODEL.eval()
        # with the test data
        test_pred = MODEL(TEST_DATA)
        # test loss
        test_loss = loss_func(test_pred.float(), TEST_LABELS)
        test_history[epoch] = test_loss.item()
        # debug
        if epoch % 10 == 0:
            test_text = f"\nTESTING: EPOCH: {epoch} TRAIN-LOSS: {loss.item()}_
→TEST-LOSS: {test_loss.item()}"
            print(colored(test_text, 'red', 'on_white'))
        elif epoch % 10 == 0:
            train_text = f"\nTRAINING: EPOCH: {epoch} LOSS: {loss.item()}"
            print(colored(train_text, 'green'))
        # zeroing
        opt.zero_grad()
        # backprop
        loss.backward()
        # stepping
        opt.step()

return MODEL, train_history, test_history

```

```

[ ]: # setting the device
device=torch.device("cuda") if torch.cuda.is_available() else torch.
→device("cpu")

# insantiate model
model = nCOV_LSTM(n_features=1, n_hidden=256,seq_len=SEQ_LENGTH, DEVICE=device,
→n_layers=3)
# sending to device
model = model.to(device)

```

```

[ ]: model, train_history, test_history = train_LSTM_model(model, X_train,
→y_train,30, device, X_test,y_test)

```

```

TESTING: EPOCH: 0 TRAIN-LOSS: 75.61992645263672 TEST-LOSS:113.64680480957031
TESTING: EPOCH: 10 TRAIN-LOSS: 45.704708099365234 TEST-LOSS:100.82435607910156
TESTING: EPOCH: 20 TRAIN-LOSS: 21.806142807006836 TEST-LOSS:85.08820343017578


```

```

[ ]: # lets look at the predictions
plt.plot(train_history, label='Training Loss')
plt.plot(test_history, label='Testing Loss')

```

```
#plt.ylim((1,10))
#plt.xlim((5,60))
plt.legend()
plt.title("Graph 3. Training vs. Testing Loss")
plt.show()
```



Predicting Daily Cases We are only predicting a single value, but we can start to predict more values by taking the predicted value (y_{pred}) as input to the LSTM and use that to make even more predictions!

```
[ ]: # lets predict more values
with torch.no_grad():
    # a single value
    test_seq = X_test[:1]
    test_seq = test_seq.to(device)
    preds = []
    for _ in range(len(X_test)):
        # predicting on a single value
        y_pred = model(test_seq)
        pred = torch.flatten(y_pred).item()
        preds.append(pred)
        #making the new sequences
        new_seqs = test_seq.cpu().numpy().flatten()
        # adding the prediction to the new sequence
        new_seqs = np.append(new_seqs, [pred])
        # using the new sequence
```



```

new_seqs = new_seqs[1:]
test_seq = torch.as_tensor(new_seqs).view(1, SEQ_LENGTH, 1).float().
→to(device)

```

```

[ ]: # now lets reverse the scaling to be able to plot it
true_cases = scaler.inverse_transform(
    np.expand_dims(y_test.flatten().numpy(), axis=0)).flatten()

# now we use the predicted cases
pred_cases = scaler.inverse_transform(np.expand_dims(preds, axis=0)).flatten()

```

```


[ ]: # lets plot to see how well we did
plt.plot(
    daily_cases.index[:len(train_data)],
    scaler.inverse_transform(train_data).flatten(),
    label='Historical Daily Cases'
)

# the other plot --> REAL
plt.plot(
    daily_cases.index[len(train_data):len(train_data)+len(true_cases)],
    true_cases,
    label='Real Daily Cases'
)

# the last plot --> PREDICTED
plt.plot(
    daily_cases.index[len(train_data):len(train_data)+len(true_cases)],
    pred_cases,
    label='Predicted Daily Cases'
)

plt.legend()
plt.show()

```



Using all the data to train the model

```
[ ]: # instantiante scaler
scaler = MinMaxScaler()
# fitting
scaler = scaler.fit(np.expand_dims(daily_cases, axis=1))
# all the data
all_data = scaler.transform(np.expand_dims(daily_cases, axis=1))
# creating the sequences
X_all, y_all = create_sequences(all_data, SEQ_LENGTH)
# to tensor
X_all = torch.from_numpy(X_all).float()
y_all = torch.from_numpy(y_all).float()
# instantiate the model
# insantiate model
x_model = nCOV_LSTM(n_features=1, n_hidden=512, seq_len=SEQ_LENGTH, ↵
↵DEVICE=device, n_layers=2)
# sending to device
x_model = x_model.to(device)
#ttraining the model
x_model, train_history, _ = train_LSTM_model(x_model, X_train, y_train, 100, ↵
↵device)
```

TRAINING: EPOCH: 0 LOSS: 47.04938507080078

TRAINING: EPOCH: 10 LOSS: 16.539140701293945

TRAINING: EPOCH: 20 LOSS: 16.53706169128418

TRAINING: EPOCH: 30 LOSS: 16.485414505004883

TRAINING: EPOCH: 40 LOSS: 16.46211814880371

TRAINING: EPOCH: 50 LOSS: 16.377498626708984

TRAINING: EPOCH: 60 LOSS: 16.325340270996094

TRAINING: EPOCH: 70 LOSS: 16.30404281616211

TRAINING: EPOCH: 80 LOSS: 16.29227066040039

TRAINING: EPOCH: 90 LOSS: 16.15147590637207

Lets predict into the future!

```
[ ]: TO_PREDICT = 50

with torch.no_grad():
    test_seq = X_all[:1]
    test_seq = test_seq.to(device)
    preds = []
    for _ in range(TO_PREDICT):
        #predicting
        y_pred = x_model(test_seq)
        pred = torch.flatten(y_pred).item()
        preds.append(pred)
        new_seq = test_seq.cpu().numpy().flatten()
        new_seq = np.append(new_seq, [pred])
        new_seq = new_seq[1:]
        test_seq = torch.as_tensor(new_seq).view(1, SEQ_LENGTH, 1).float().
        ↪to(device)

# the predicted cases
pred_cases = scaler.inverse_transform(
    np.expand_dims(preds, axis=0)
).flatten()
```


```

daily_cases.index[-1]

# we create a predicted index range
pred_idx = pd.date_range(
    start=daily_cases.index[-1],
    periods=TO_PREDICT+1,
    closed='right'
)
# get a series out of it
pred_cases = pd.Series(
    data = pred_cases,
    index=pred_idx
)

#plotting
plt.plot(
    pred_cases,
    label='Predicted Daily Cases'
)
plt.legend()
plt.title("Graph 4. Predicted Daily Cases Using all the data")
plt.show()

```



```

[ ]: # using all the data
plt.plot(daily_cases, label='Historical Daily Cases')

```

```
plt.plot(pred_cases , label='Predicted Daily Cases' )  
plt.legend()  
plt.title("Graph 5. Real + Prediction Graph ")  
plt.show()
```

