

# Chapter\_5\_LSTM\_Autoencoder

January 28, 2022

## 1 Chapter 5: LSTM Anomaly Detection with Autoencoders

### Imports

```
[ ]: !pip install -qq arff2pandas
```

```
Building wheel for arff2pandas (setup.py) ... done
Building wheel for liac-arff (setup.py) ... done
```

```
[ ]: import os
      #linalg
      import pandas as pd
      import numpy as np
      #plotting
      import matplotlib.pyplot as plt
      import seaborn as sns
      from pylab import rcParams

      # split
      from sklearn.model_selection import train_test_split

      #loading arff files
      from arff2pandas import a2p

      # torch
      import torch
      from torch import nn, optim
      import torch.nn.functional as F

      #copy
      import copy
      #colored print
      import termcolor
      from termcolor import colored

      # set the size of the figures
      rcParams['figure.figsize'] = 15, 8
```

```
# setting the device
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
[ ]: device
```

```
[ ]: device(type='cuda')
```

### Fetch data

```
[ ]: !wget --id 16MIleqoIr1vYxlGk4GKnGmrsCPuWkkpT
```

Downloading...

From: <https://drive.google.com/uc?id=16MIleqoIr1vYxlGk4GKnGmrsCPuWkkpT>

To: /content/ECG5000.zip

100% 10.6M/10.6M [00:00<00:00, 64.9MB/s]

```
[ ]: !unzip -qq ECG5000.zip
```

### Load Data

```
[ ]: with open('ECG5000_TRAIN.arff') as f:
    train = a2p.load(f)

    with open('ECG5000_TEST.arff') as k:
        test = a2p.load(k)

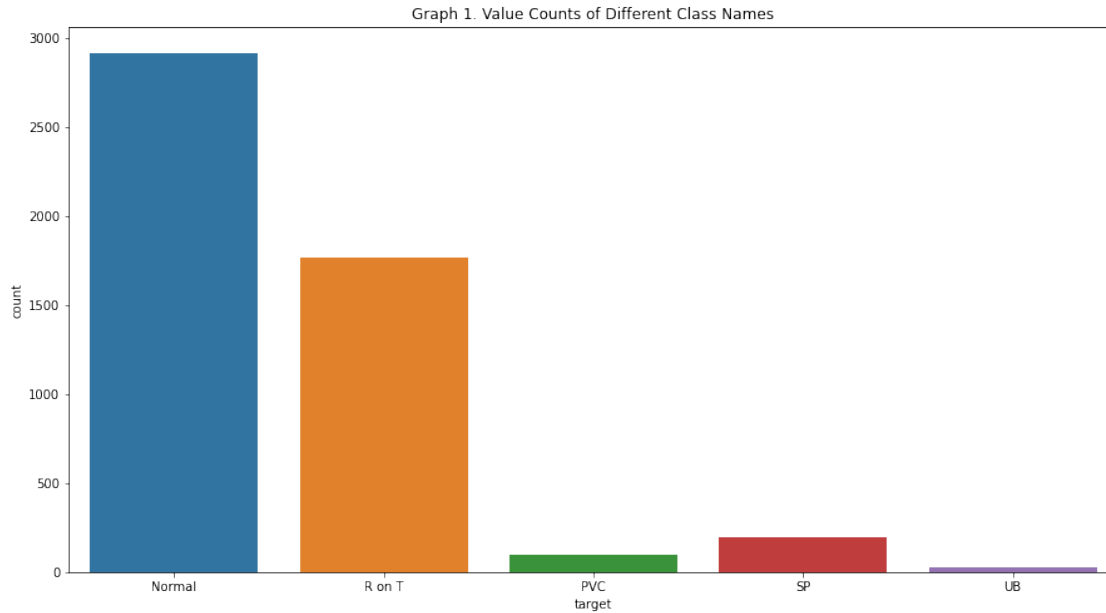
    # into a single dataframe
    df = train.append(test)
    df = df.sample(frac=1.0)
    # rename the last column
    cols = list(df.columns)
    cols[-1] = 'target'
    df.columns = cols

    # change the dtype of the target column
    df['target'] = df['target'].astype(int)
```

### Visualize Data

```
[ ]: # some class names
NORMAL_CLASS = 1
class_names = ['Normal', 'R on T', 'PVC', 'SP', 'UB']

# plotting
ax = sns.countplot(x=df.target)
ax.set_xticklabels(class_names)
plt.title('Graph 1. Value Counts of Different Class Names')
plt.show()
```



```
[ ]: # from the google colab from the book
def plot_time_series_class(data, class_name, ax, n_steps=10):
    time_series_df = pd.DataFrame(data)

    smooth_path = time_series_df.rolling(n_steps).mean()
    path_deviation = 2 * time_series_df.rolling(n_steps).std()

    under_line = (smooth_path - path_deviation)[0]
    over_line = (smooth_path + path_deviation)[0]

    ax.plot(smooth_path, linewidth=2)
    ax.fill_between(
        path_deviation.index,
        under_line,
        over_line,
        alpha=.125
    )
    ax.set_title(class_name)

classes = df.target.unique()

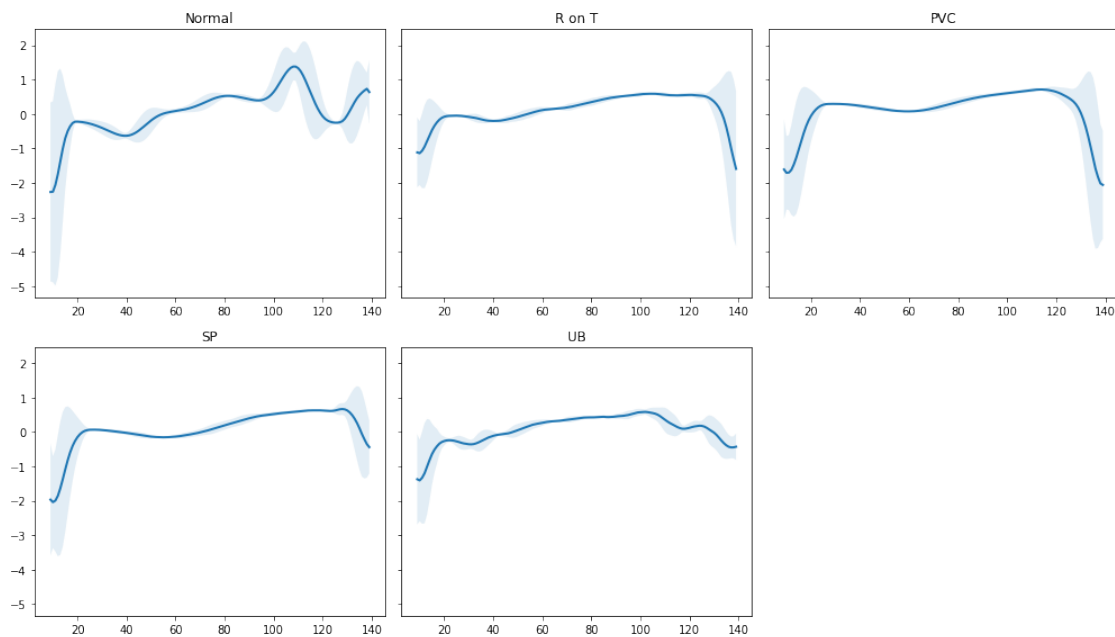
fig, axs = plt.subplots(
    nrows=len(classes) // 3 + 1,
    ncols=3,
    sharey=True,
    figsize=(14, 8)
)
```

```

for i, cls in enumerate(classes):
    ax = axs.flat[i]
    data = df[df.target == cls] \
        .drop(labels='target', axis=1) \
        .mean(axis=0) \
        .to_numpy()
    plot_time_series_class(data, class_names[i], ax)

fig.delaxes(axs.flat[-1])
fig.tight_layout();
plt.show()

```



## Separating Datasets

```

[ ]: # lets separate the data: Normal & ~Normal
normal_df = df[df['target'] == NORMAL_CLASS].drop('target', axis=1)
# ~Normal
abnormal_df = df[df['target'] != NORMAL_CLASS].drop('target', axis=1)
# showing the shapes
print("Normal and abnormal dataframes")
print(normal_df.shape, abnormal_df.shape)

# lets split the data
# train and validation
# 0.15 * 0.33 5%, thus: splits = train:85%, val:10%, test:5%

```

```
df_train, df_val = train_test_split(normal_df, test_size=0.15, random_state=69)
# val and testing
df_val, df_test = train_test_split(df_val, test_size=0.33, random_state=69)
```

Normal and abnormal dataframes  
(2919, 140) (2081, 140)

## Making Datasets

```
[ ]: # now we need to create the datasets
# we need to convert it into shape:
# (sequence_length * n_features)
def make_dataset(DF, mode='train'):
    # first we need to flatten the dataframe and into a list
    seqs = DF.astype(np.float32).to_numpy().tolist()
    # making the dataset: converting to tensors
    # we add a dimension with unsqueeze at the 2nd position --> (141,1)
    ds = [torch.tensor(s).unsqueeze(1).float() for s in seqs]
    # we then stack these, and get their shape
    # stacking is done similar to hstack or vstack in numpy
    n_seq, seq_len, n_features = torch.stack(ds).shape
    if mode == 'train':
        return ds, seq_len, n_features
    else:
        return ds

# we make the datasets
# TRAINING
train_dataset, seq_len, n_features = make_dataset(df_train, mode='train')
# VALIDATION
val_dataset = make_dataset(df_val, mode='not-train')
# TESTING
test_dataset = make_dataset(df_test, mode='not-train')
# ABNORMAL DATASET
test_abnormal_ds = make_dataset(abnormal_df, mode='not-normal')
```

## Making the Encoder

```
[ ]: # ENCODER
class Encoder(nn.Module):
    def __init__(self, seq_len, n_features, embedding_dim=64):
        super(Encoder, self).__init__()
        #selfs
        self.seq_len = seq_len
        self.n_features = n_features
        self.embedding_dim = embedding_dim
        # It is twice the embedding size because we want to
        # expand it so it can have more dimensions
```

```

# therefore there is more dimensions / features that can be
# compressed into latent variables
self.hidden_dim = 2 * embedding_dim

# making the recurrent network
self.rnn1 = nn.LSTM(
    input_size = n_features,
    hidden_size = self.hidden_dim,
    num_layers=1,
    batch_first=True
)

# the second recurrent network
self.rnn2 = nn.LSTM(
    input_size = self.hidden_dim,
    hidden_size = embedding_dim,
    num_layers=1,
    batch_first=True
)

# forward pass
def forward(self, x):
    x = x.reshape((1, self.seq_len, self.n_features))
    x, (_, _) = self.rnn1(x)
    x, (hidden_n, _) = self.rnn2(x)
    return hidden_n.reshape((self.n_features, self.embedding_dim))

```

## Making the Decoder

```

[ ]: class Decoder(nn.Module):
    def __init__(self, seq_len, input_dim=64, n_features=1):
        super(Decoder, self).__init__()
        self.seq_len = seq_len
        self.input_dim = input_dim
        self.hidden_dim = 2 * input_dim
        self.n_features = n_features

        # making the recurrent network
        self.rnn1 = nn.LSTM(
            input_size = input_dim,
            hidden_size = input_dim,
            num_layers=1,
            batch_first=True
        )

        # the second recurrent network
        self.rnn2 = nn.LSTM(
            input_size = input_dim,
            hidden_size = self.hidden_dim,
            num_layers=1,

```

```

        batch_first=True
    )

    # output layer
    self.output_layer = nn.Linear(self.hidden_dim, n_features)

    #forward pass
    def forward(self, x):
        x = x.repeat(self.seq_len, self.n_features)
        x = x.reshape((self.n_features, self.seq_len, self.input_dim))
        x, (hidden_n, cell_n) = self.rnn1(x)
        x, (hidden_n, cell_n) = self.rnn2(x)
        x = x.reshape((self.seq_len, self.hidden_dim))
        return self.output_layer(x)

```

## Making the Autoencoder (Encoder + Decoder)

```

[ ]: class RAE(nn.Module):
    def __init__(self, seq_len, n_features, embedding_dim=64):
        super(RAE, self).__init__()
        self.encoder = Encoder(seq_len, n_features, embedding_dim).to(device)
        self.decoder = Decoder(seq_len, embedding_dim, n_features).to(device)

    #forwards pass
    def forward(self, x):
        x = self.encoder(x)
        x = self.decoder(x)
        return x

```

## Instantiate the model

```

[ ]: model_ae = RAE(seq_len, n_features, 128)
    model_ae = model_ae.to(device)

```

```

[ ]: model_ae

```

```

[ ]: RAE(
  (encoder): Encoder(
    (rnn1): LSTM(1, 256, batch_first=True)
    (rnn2): LSTM(256, 128, batch_first=True)
  )
  (decoder): Decoder(
    (rnn1): LSTM(128, 128, batch_first=True)
    (rnn2): LSTM(128, 256, batch_first=True)
    (output_layer): Linear(in_features=256, out_features=1, bias=True)
  )
)

```

## Training Loop

```
[ ]: def training_loop(MODEL, TRAIN_DATA, VAL_DATA, DEVICE, EPOCHS):  
    '''  
    Helper function to train the model on the:  
    - Train dataset  
    - Validation dataset  
    for a specified number of EPOCHS  
    '''  
  
    # define the optimizer  
    opt = optim.Adam(MODEL.parameters(), lr=1e-3)  
    # define our cost function  
    loss_func = nn.L1Loss(reduction='sum').to(DEVICE)  
    # track the history  
    history = dict(train=[], val=[])  
  
    # state dict  
    #best_model = copy.deepcopy(MODEL.state_dict())  
    best_loss = 10000.0  
  
    # loop  
    for epoch in range(EPOCHS):  
        # set to training  
        MODEL = MODEL.train()  
        # track the loss  
        train_loss = []  
        for original_sequence in TRAIN_DATA:  
            # zero the gradients  
            opt.zero_grad()  
            # send to device  
            original_sequence = original_sequence.to(DEVICE)  
            # predicting the sequence  
            pred_seq = MODEL(original_sequence)  
            # determine the loss  
            loss = loss_func(pred_seq, original_sequence)  
            # propagate  
            loss.backward()  
            # step  
            opt.step()  
            # track the loss  
            train_loss.append(loss.item())  
        # NOW WE VALIDATE  
        val_losses = []  
        MODEL = MODEL.eval()  
        with torch.no_grad():  
            for original_sequence in VAL_DATA:  
                # send to d
```



```

original_sequence = original_sequence.to(DEVICE)
#predict
pred_seq = MODEL(original_sequence)
# calculate the loss
val_loss = loss_func(pred_seq, original_sequence)
# append
val_losses.append(val_loss.item())
# taking the average
train_loss = np.mean(train_loss)
val_losses = np.mean(val_losses)
# appending to the history
history['train'].append(train_loss)
history['val'].append(val_losses)
# now keep track of the best loss
if val_loss < best_loss:
    best_loss = val_loss
    #best_model = copy.deepcopy(MODEL.state_dict())
    debug_txt = f"\nTRAINING --- EPOCH: {epoch} --- TRAIN LOSS: {train_loss}┐
↪--- VAL LOSS: {val_losses}"
    print(colored(debug_txt, 'red', 'on_white'))
# loading the best model
#MODEL.load_state_dict(best_model)
return MODEL.eval(), history

```

## Actual Training

```

[ ]: model, history = training_loop(model_ae,
                                   train_dataset,
                                   val_dataset,
                                   device,
                                   EPOCHS=30)

```

```

TRAINING --- EPOCH: 0 --- TRAIN LOSS: 20.565458772452498 --- VAL LOSS:18.93135979558014
TRAINING --- EPOCH: 1 --- TRAIN LOSS: 19.430236712428457 --- VAL LOSS:18.16281157548924
TRAINING --- EPOCH: 2 --- TRAIN LOSS: 18.784084511110734 --- VAL LOSS:19.45093519126189
TRAINING --- EPOCH: 3 --- TRAIN LOSS: 18.59398806935979 --- VAL LOSS:16.153816123871266
TRAINING --- EPOCH: 4 --- TRAIN LOSS: 17.97936878834747 --- VAL LOSS:15.848345710962706
TRAINING --- EPOCH: 5 --- TRAIN LOSS: 17.473653324185626 --- VAL LOSS:16.302490787701394
TRAINING --- EPOCH: 6 --- TRAIN LOSS: 18.100933074374584 --- VAL LOSS:21.889322668212266
TRAINING --- EPOCH: 7 --- TRAIN LOSS: 18.19583534243029 --- VAL LOSS:15.41353888859668
TRAINING --- EPOCH: 8 --- TRAIN LOSS: 17.531168819482843 --- VAL LOSS:14.892759201062825
TRAINING --- EPOCH: 9 --- TRAIN LOSS: 16.898734957015787 --- VAL LOSS:21.670208556660207
TRAINING --- EPOCH: 10 --- TRAIN LOSS: 16.732507542514455 --- VAL LOSS:14.960598153058987
TRAINING --- EPOCH: 11 --- TRAIN LOSS: 15.98129130928135 --- VAL LOSS:18.521030761276087
TRAINING --- EPOCH: 12 --- TRAIN LOSS: 15.784048732948227 --- VAL LOSS:15.092466087471504
TRAINING --- EPOCH: 13 --- TRAIN LOSS: 15.637804906630603 --- VAL LOSS:14.3868973442312
TRAINING --- EPOCH: 14 --- TRAIN LOSS: 15.467676572288443 --- VAL LOSS:14.831088966070181
TRAINING --- EPOCH: 15 --- TRAIN LOSS: 15.403793454698764 --- VAL LOSS:13.585212720538976

```

```

TRAINING --- EPOCH: 16 --- TRAIN LOSS: 15.409222868458487 --- VAL LOSS:13.27314734703038
TRAINING --- EPOCH: 17 --- TRAIN LOSS: 14.838274210415172 --- VAL LOSS:21.413372306693535
TRAINING --- EPOCH: 18 --- TRAIN LOSS: 14.676855360005758 --- VAL LOSS:14.445762618003037
TRAINING --- EPOCH: 19 --- TRAIN LOSS: 14.521285603864209 --- VAL LOSS:14.29533950786135
TRAINING --- EPOCH: 20 --- TRAIN LOSS: 14.78504458938285 --- VAL LOSS:23.175373607934944
TRAINING --- EPOCH: 21 --- TRAIN LOSS: 14.71017585632927 --- VAL LOSS:13.903403591377337
TRAINING --- EPOCH: 22 --- TRAIN LOSS: 14.053567840417802 --- VAL LOSS:16.598857736424783
TRAINING --- EPOCH: 23 --- TRAIN LOSS: 14.369419816711746 --- VAL LOSS:13.837991559871634
TRAINING --- EPOCH: 24 --- TRAIN LOSS: 14.103604733583573 --- VAL LOSS:14.812460271164825
TRAINING --- EPOCH: 25 --- TRAIN LOSS: 13.658300871620348 --- VAL LOSS:13.194585788778884
TRAINING --- EPOCH: 26 --- TRAIN LOSS: 13.871981800683608 --- VAL LOSS:12.956847579812841
TRAINING --- EPOCH: 27 --- TRAIN LOSS: 13.680341306591842 --- VAL LOSS:13.655689110934938
TRAINING --- EPOCH: 28 --- TRAIN LOSS: 13.576541122627182 --- VAL LOSS:16.008496596951534
TRAINING --- EPOCH: 29 --- TRAIN LOSS: 13.504578802000367 --- VAL LOSS:16.200924323280518

```

### Defining a predict function

```

[ ]: # create a function to predict
def _predict(MODEL, DATASET):
    predictions, losses = [],[]
    loss_func = nn.L1Loss(reduction='sum').to(device)
    # test mode
    with torch.no_grad():
        MODEL = MODEL.eval()
        for original_sequence in DATASET:
            # to d
            original_sequence = original_sequence.to(device)
            # pred
            pred_seq = MODEL(original_sequence)
            # loss
            loss = loss_func(pred_seq, original_sequence)
            # track
            predictions.append(pred_seq.cpu().detach().numpy().flatten())
            losses.append(loss.item())
    return predictions, losses

```

### Evaluation

```

[ ]: # count the number of correct predictions
def get_correct(VALUE, LOSSES, DATASET,mode='Normal',show_graph=True):
    if show_graph:
        g = sns.distplot(LOSSES, bins=50, kde=True)
        plt.show()

    if mode == 'Normal':
        correct = sum(k<= VALUE for k in LOSSES)
        text = f"\nCorrect Normal Predictions --- {correct}/{len(DATASET)}"
        print(colored(text, 'green','on_white',attrs=['bold']))

```

```

if mode == 'Abnormal':
    correct = sum(uu > VALUE for uu in LOSSES)
    text = f"\nCorrect Abnormal Predictions --- {correct}/{len(DATASET)}"
    print(colored(text, 'red','on_white',attrs=['bold']))

```

```

[ ]: # lets define a threshold
THRESHOLD = 26 # defines what is normal vs abnormal

# lets look at the normal heartbeats
predictions, pred_losses = _predict(model_ae, test_dataset)

# abnormal ones
# filter the abnormal dataset to match length of original
anomaly_dataset = test_abnormal_ds[:len(test_dataset)]
ab_preds, ab_losses = _predict(model_ae,anomaly_dataset)

```


```

[ ]: # looking at the normal graphs
get_correct(THRESHOLD, pred_losses, test_dataset ,mode='Normal',show_graph=True)

```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)




```
Correct Normal Predictions --- 137/145
```

```
[ ]: #looking at the abnormal
get_correct(THRESHOLD, ab_losses,anomaly_dataset,
↳mode='Abnormal',show_graph=True)
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619:  
FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

```
warnings.warn(msg, FutureWarning)
```



Correct Abnormal Predictions --- 144/145

Looking at more examples

```
[ ]: def plot_prediction(data, model, title, ax):
    predictions, pred_losses = _predict(model, [data])

    ax.plot(data, label='true')
    ax.plot(predictions[0], label='reconstructed')
    ax.set_title(f'{title} (loss: {np.around(pred_losses[0], 2)})')
    ax.legend()

fig, axs = plt.subplots(
    nrows=2,
    ncols=6,
    sharey=True,
    sharex=True,
    figsize=(22, 8))
```

```


)

for i, data in enumerate(test_dataset[:6]):
    plot_prediction(data, model_ae, title='Normal', ax=axes[0, i])

for i, data in enumerate(test_abnormal_ds[:6]):
    plot_prediction(data, model_ae, title='Anomaly', ax=axes[1, i])

fig.tight_layout();

```



[ ]: